



EDUCATION

Carnegie Mellon University

Ph.D. in Computer Science

2017 - 2022 (Expected) | PA, USA

Advised by Todd C. Mowry

Carnegie Mellon University

M.Sc. in Computer Science

2015 - 2016 | PA, USA

Major GPA: 3.9 / 4.0

Simon Fraser University

B.Sc. in Computer Science

2012 - 2015 | BC, Canada

Graduated with Distinction

Zhejiang University

B.Eng. in Computer Science

2010 - 2015 | Zhejiang, China

Major GPA: 4.46 / 5.0 (Avg. 89.2 / 100)

INTERESTS

Computer Architecture

Memory & Cache Subsystem Optimization

Database Concurrency Control

Concurrent Data Structures

AWARDS

Open Scholarship, SFU (Spring 2015)

Open Scholarship, SFU (Fall 2014)

Dean's Honor Roll, SFU (Spring 2014)

President's Honor Roll, SFU (Fall 2013)

Open Scholarship, SFU (Spring 2013)

Entrance Scholarship, SFU (Fall 2012)

National Scholarship of ZU (Fall 2012)

Department Scholarship of ZU (Fall 2011)

LINKS

Github

<https://github.com/wangziqi2013/>

<https://github.com/wangziqi2016/>

Homepage / Blog

<https://wangziqi2013.github.io/>

COMMUNITY

Contributed to [stamp-rtm](#) and [synchrobench](#)

Open-sourced Microsoft Research's BwTree as [OpenBwTree](#)

DOCTORAL RESEARCH

Cache Compression on 2D Address Space

2020 - Present | Supervised by Todd Mowry and Michael Kozuch

- Extends physical address space with a second dimension, the version dimension, by tagging cache lines in the hierarchy with version numbers.
- Software provides compression hints by aligning homogeneous cache lines (lines with identical layout but slightly different data) to the same address but on different versions.
- Hardware delta-encodes aligned cache lines for better compressability.
- Supports per-page, dynamic configuration of compressed block layouts, avoiding penalizing applications for disabling compression and retaining backward compatibility.

High-Frequency Full System Persistence with NVDIMM

2018 - 2020 | Supervised by Todd Mowry and Michael Kozuch

- Hardware-assisted, high-frequency address space checkpointing for supporting durable transactions and data structures, system backups, crash recovery, etc.
- Features lower write amplification than logging-based approaches. Cycle overhead is comparable to or lower than other NVM-based checkpointing designs. Demonstrates better scalability.
- Implemented with coherence-driven dirty data tracking and shadow-mapped NVM store. Fully distributed with Lamport clock algorithm.
- Modeled with zSim. Results show 29% - 47% less writes and memory bandwidth consumption than previous state-of-the-art.

Main Memory Multiversioning and Hardware Transactional Memory with Reduced Abort Rates

2017 - 2018 | Supervised by Todd Mowry and Michael Kozuch

- Extends existing coherence protocol to support multiversioning on the physical address space. Addresses are tagged with an extra version number and accessed by hardware version lookup.
- Proposes OverlayTM, a Hardware Transactional Memory design featuring arbitrarily large transactions and fast read-only commit. Speculative cache lines are managed as hardware versions. Conflicts are detected with Optimistic Concurrency Control.
- Modeled with zSim. Results show significant reduction in aborts and less wasted cycles than Intel TSX.

MASTERS RESEARCH

Efficient In-Memory Indexing with Lock-Free B+ Tree

2016 - 2017 | Supervised by Andrew Pavlo

- Worked on Peloton (renamed to terrier in 2018), a self-driving database management system. Responsible for implementing, profiling and tuning of the index subsystem.
- Implemented a highly optimized BwTree, a concurrent lock-free B+Tree index proposed by Microsoft Research. Currently deployed as Peloton's in-memory index.
- Benchmarked BwTree under YCSB workloads. Analyzed BwTree performance characteristics and compared it with Optimistic Lock Coupling and Masstree.
- Open-sourced the optimized BwTree as OpenBwTree and the index benchmarking framework.

INTERDISCIPLINARY RESEARCH

Development of Flash Translation Layer Simulator for Grading

Jan 2016 – Apr 2016 | Independent Study | Supervised by Greg Ganger

- Developed a flash simulator for 15-746 Fall 2016 as the course project evaluation platform.

Analyzing a Million Apps from Android Store for Privacy Leakage

Jan 2017 – Jul 2017 | Research Scientist | Supervised by Sebastian Zimmeck

Static Analysis of Android Apps for Privacy Leakage

Sep 2015 – Dec 2015 | Research Assistant | Supervised by Norman Sadeh

- Developed a Python framework for static data-flow analysis of Android apps.
- Scaled the framework to be running on Amazon cloud platform for million-scale app analysis.

PUBLICATION

2019

Wang, Ziqi, Vivek Seshadri, Todd C. Mowry, and Michael Kozuch. "Multiversed Page Overlays: Enabling Faster Serializable Hardware Transactional Memory." In 2019 28th International Conference on Parallel Architectures and Compilation Techniques (PACT). IEEE, 2019.

Zimmeck, Sebastian, Peter Story, Daniel Smullen, Abhilasha Ravichander, Ziqi Wang, Joel Reidenberg, N. Cameron Russell, and Norman Sadeh. "MAPS: Scaling privacy compliance analysis to a million apps." Proceedings on Privacy Enhancing Technologies 2019, no. 3 (2019): 66-86.

2018

Wang, Ziqi, Andrew Pavlo, Hyeontaek Lim, Viktor Leis, Huanchen Zhang, Michael Kaminsky, and David G. Andersen. "Building a bw-tree takes more than just buzz words." In Proceedings of the 2018 International Conference on Management of Data, pp. 473-488. ACM, 2018.

2017

Pavlo, Andrew, Gustavo Angulo, Joy Arulraj, Haibin Lin, Jiexi Lin, Lin Ma, Prashanth Menon et al. "Self-Driving Database Management Systems." In CIDR, vol. 4, p. 1. 2017. (Authors sorted by last name)

2016

Zimmeck, Sebastian, Ziqi Wang, Lieyong Zou, Roger Iyengar, Bin Liu, Florian Schaub, Shomir Wilson, Norman Sadeh, Steven Bellovin, and Joel Reidenberg. "Automated analysis of privacy requirements for mobile apps." In 2016 AAAI Fall Symposium Series. 2016.

TEACHING

Parallel Computer Architecture and Programming

CMU 15-418/618 | Fall 2020 | with Todd Mowry

Graduate Computer Architectures

CMU 15-740 | Fall 2018 | with Nathan Backmann

Storage Systems

CMU 15-746 | Fall 2016 | with Greg Ganger and Garth Gibson

Distributed Systems

CMU 15-440 | Spring 2016 | with Mahadev Satyanarayanan

INDUSTRY

Ericsson Canada | Software Engineer Co-op

May 2014 - Dec 2014 | Burnaby, BC, Canada

- Added new feature to Ericsson's proprietary IP Operating System (IPOS) packet classifier
- Developed a Ternary Content Addressable Memory (TCAM) image rebuilder
- Developed a testing framework for PPPoE on routers

RESEARCH TALKS

Concurrency Control Methods for Hardware Transactions

Parallel reading group, Spring 2019

Hardware Transactional Memory

Guest speaker in 15-740 Computer Architecture, 2018 Fall

Building a Bw-Tree Takes More Than Just Buzz Words

SIGMOD 2018 Paper Presentation

Efficient Memory Indexing and Lock-Freedom

Oracle Research, 2017